

Introduction

Final Project Submission

- Student Name: Adam Marianacci
- Student Pace: Flex
- Scheduled project review date/time: TBD
- Instructor Name: Mark Barbour

Business Understanding

It is my job to help the WWFA (Water Wells For Africa) locate wells that need to be repaired in Tanzania.

Data Understanding

Data Preperation

```
In [1]: 1 # Importing the necessary libraries
2 import pandas as pd
3 from datetime import datetime
4 import numpy as np
5 import seaborn as sns
6 import statsmodels as sm
7 import sklearn
8 import sklearn.preprocessing as preprocessing
9 import matplotlib.pyplot as plt
10 from scipy import stats
11 from sklearn import linear_model
12 from sklearn.linear_model import LogisticRegression
13 from sklearn.feature_selection import RFE
14 from sklearn.tree import DecisionTreeClassifier
15 from sklearn.model_selection import train_test_split
16 from sklearn.preprocessing import MinMaxScaler
17 from sklearn.linear_model import LinearRegression
18 from sklearn.preprocessing import OneHotEncoder
19 from sklearn.compose import ColumnTransformer
20 from sklearn.impute import SimpleImputer
21 from sklearn.metrics import r2_score, mean_squared_error, mean_absolute_error
22 import warnings
23 warnings.filterwarnings('ignore')
```

```
In [2]: 1 # Set display options to show all rows and columns
2 pd.set_option('display.max_rows', None)
3 pd.set_option('display.max_columns', None)
```

```
In [3]: 1 # Importing the dataframes
2 df_x = pd.read_csv('data/training_set_values.csv')
3 df_y = pd.read_csv('data/training_set_labels.csv')
```

```
In [4]: 1 # Combining the 2 dataframes into 1 new dataframe
2 Waterwells_df = pd.concat([df_y, df_x], axis=1)
```

```
In [5]: 1 # Previewing the dataframe
        2 Waterwells_df.head()
```

Out[5]:

	id	status_group	id	amount_tsh	date_recorded	funder	gps_height	installer	longitude	latitude	wpt_name	num_priv
0	69572	functional	69572	6000.0	2011-03-14	Roman	1390	Roman	34.938093	-9.856322	none	
1	8776	functional	8776	0.0	2013-03-06	Grumeti	1399	GRUMETI	34.698766	-2.147466	Zahanati	
2	34310	functional	34310	25.0	2013-02-25	Lottery Club	686	World vision	37.460664	-3.821329	Kwa Mahundi	
3	67743	non functional	67743	0.0	2013-01-28	Unicef	263	UNICEF	38.486161	-11.155298	Zahanati Ya Nanyumbu	
4	19728	functional	19728	0.0	2011-07-13	Action In A	0	Artisan	31.130847	-1.825359	Shuleni	

```
In [6]: 1 #dropping the missing values from the 'construction_year' column and creating a new df
        2 Construction_Year_df = Waterwells_df[Waterwells_df['construction_year'] != 0]
        3
        4 # Calculate the current year
        5 current_year = datetime.now().year
        6
        7 # Create a new column 'age' by subtracting construction year from the current year
        8 Construction_Year_df['age'] = current_year - Waterwells_df['construction_year']
```

```
In [7]: 1 Construction_Year_df['construction_year'].value_counts()
```

```
Out[7]: 2010      2645
        2008      2613
        2009      2533
        2000      2091
        2007      1587
        2006      1471
        2003      1286
        2011      1256
        2004      1123
        2012      1084
        2002      1075
        1978      1037
        1995      1014
        2005      1011
        1999       979
        1998       966
        1990       954
        1985       945
        1996       811
        1980       811
        1984       779
        1982       744
        1994       738
        1972       708
        1974       676
        1997       644
        1992       640
        1993       608
        2001       540
        1988       521
        1983       488
        1975       437
        1986       434
        1976       414
        1970       411
        1991       324
        1989       316
        1987       302
        1981       238
```

1977	202
1979	192
1973	184
2013	176
1971	145
1960	102
1967	88
1963	85
1968	77
1969	59
1964	40
1962	30
1961	21
1965	19
1966	17

Name: construction_year, dtype: int64

Dropping columns that are not directly related to the business problem and also have high cardinality, making them difficult to one hot encode.

```

In [8]: 1 # Dropping irrelevant columns from the dataframe, also columns with large amounts of missing data
        2 columns_to_drop = [
        3     'id', 'scheme_management', 'region', 'payment', 'public_meeting', 'district_code', 'population',
        4     'num_private', 'basin', 'construction_year',
        5     'waterpoint_type_group', 'source_class', 'payment_type', 'management_group', 'recorded_by',
        6     'extraction_type', 'management',
        7     'source_type', 'extraction_type_group', 'permit', 'funder',
        8     'date_recorded', 'installer', 'ward', 'scheme_name', 'wpt_name', 'lga', 'subvillage'
        9 ]
        10
        11 Waterwells_df = Waterwells_df.drop(columns_to_drop, axis=1, errors='ignore')
        12
  
```

```
In [9]: 1 # Examining the dimensions of the dataframe
        2 Waterwells_df.head()
```

Out[9]:

	status_group	amount_tsh	gps_height	longitude	latitude	region_code	extraction_type_class	water_quality	quality_group	quantity
0	functional	6000.0	1390	34.938093	-9.856322	11	gravity	soft	good	enough
1	functional	0.0	1399	34.698766	-2.147466	20	gravity	soft	good	insufficient
2	functional	25.0	686	37.460664	-3.821329	21	gravity	soft	good	enough
3	non functional	0.0	263	38.486161	-11.155298	90	submersible	soft	good	enough
4	functional	0.0	0	31.130847	-1.825359	18	gravity	soft	good	sea level

In [10]:

```
1 # Checking for missing values and learning about the datatypes of the columns
2 Waterwells_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 59400 entries, 0 to 59399
Data columns (total 13 columns):
#   Column                Non-Null Count  Dtype
---  -
0   status_group          59400 non-null  object
1   amount_tsh            59400 non-null  float64
2   gps_height            59400 non-null  int64
3   longitude              59400 non-null  float64
4   latitude              59400 non-null  float64
5   region_code           59400 non-null  int64
6   extraction_type_class 59400 non-null  object
7   water_quality          59400 non-null  object
8   quality_group         59400 non-null  object
9   quantity              59400 non-null  object
10  quantity_group         59400 non-null  object
11  source                 59400 non-null  object
12  waterpoint_type        59400 non-null  object
dtypes: float64(3), int64(2), object(8)
memory usage: 5.9+ MB
```

```
In [11]: 1 # Create a new column 'needs_repair' by merging the two categories
2 Waterwells_df['needs_repair'] = Waterwells_df['status_group'].replace(
3     {'functional': 0, 'non functional': 1,
4     'functional but needs repair': 1})
5
6 # Drop the original 'status_group' column
7 Waterwells_df.drop('status_group', axis=1, inplace=True)
8
9 #Display the updated DataFrame
10 Waterwells_df.head()
11
12
```

Out[11]:

	amount_tsh	gps_height	longitude	latitude	region_code	extraction_type_class	water_quality	quality_group	quantity	quantity
0	6000.0	1390	34.938093	-9.856322	11	gravity	soft	good	enough	
1	0.0	1399	34.698766	-2.147466	20	gravity	soft	good	insufficient	ins
2	25.0	686	37.460664	-3.821329	21	gravity	soft	good	enough	
3	0.0	263	38.486161	-11.155298	90	submersible	soft	good	dry	
4	0.0	0	31.130847	-1.825359	18	gravity	soft	good	seasonal	si

```
In [12]: 1 # Defining X and y variables
2 y = Waterwells_df["needs_repair"]
3 X = Waterwells_df.drop("needs_repair", axis=1)
```

```
In [13]: 1 # Performing a train, test, split
2 X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.2, random_state=42)
```



```
In [14]: 1 # Looking at the number of missing values in each column
        2 X_train.isna().sum()
```

```
Out[14]: amount_tsh          0
gps_height          0
longitude           0
latitude            0
region_code         0
extraction_type_class 0
water_quality        0
quality_group        0
quantity            0
quantity_group       0
source              0
waterpoint_type      0
dtype: int64
```

```
In [15]: 1 #Defining categorical df
        2 X_train_categorical = X_train.select_dtypes(include='object').copy()
        3 X_train_categorical.head()
        4
```

```
Out[15]:
```

	extraction_type_class	water_quality	quality_group	quantity	quantity_group	source	waterpoint_type
3607	gravity	soft	good	insufficient	insufficient	spring	communal standpipe
50870	handpump	soft	good	enough	enough	shallow well	hand pump
20413	other	soft	good	enough	enough	shallow well	other
52806	gravity	soft	good	insufficient	insufficient	river	communal standpipe
50091	other	salty	salty	enough	enough	shallow well	other

```
In [16]: 1 #Inspecting the unique values of source
        2 X_train_categorical['quality_group'].unique()
```

```
Out[16]: array(['good', 'salty', 'unknown', 'colored', 'fluoride', 'milky'],
              dtype=object)
```

```
In [17]: 1 X_train_categorical['quality_group'].value_counts()
```

```
Out[17]: good          40633  
salty           4173  
unknown         1490  
milky           650  
colored          395  
fluoride         179  
Name: quality_group, dtype: int64
```

```
In [18]: 1 # Removing 'other' and 'unknown' from the source column  
2 X_train_categorical = X_train_categorical[~X_train_categorical['source'].isin(['other', 'unknown'])]  
3  
4 # Display the updated counts  
5 print(X_train_categorical['source'].value_counts())  
6
```

```
shallow well      13540  
spring            13537  
machine dbh       8849  
river             7719  
rainwater harvesting 1829  
hand dtw          701  
lake              606  
dam               505  
Name: source, dtype: int64
```

cite this

```
In [19]: 1 ohe = OneHotEncoder(handle_unknown="ignore", sparse=False, drop='first').set_output(transform='x')
2 # column_to_encode = X_train_categorical['source']
3 X_train_categorical = ohe.fit_transform(X_train_categorical.values.reshape(-1,1))
4 X_train_categorical.head()
```

Out[19]:

	x0_colored	x0_coloured	x0_communal standpipe	x0_communal standpipe multiple	x0_dam	x0_dry	x0_enough	x0_fluoride	x0_fluoride abandoned	x0_good	x0_gravity	x
0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	1.0	
1	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	
2	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	1.0	0.0	
3	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	
4	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	

```
In [20]: 1 type(X_train_categorical)
```

Out[20]: pandas.core.frame.DataFrame

In [21]:

```
1 X_train_categorical.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 331002 entries, 0 to 331001
```

```
Data columns (total 33 columns):
```

#	Column	Non-Null Count	Dtype
0	x0_colored	331002 non-null	float64
1	x0_coloured	331002 non-null	float64
2	x0_communal standpipe	331002 non-null	float64
3	x0_communal standpipe multiple	331002 non-null	float64
4	x0_dam	331002 non-null	float64
5	x0_dry	331002 non-null	float64
6	x0_enough	331002 non-null	float64
7	x0_fluoride	331002 non-null	float64
8	x0_fluoride abandoned	331002 non-null	float64
9	x0_good	331002 non-null	float64
10	x0_gravity	331002 non-null	float64
11	x0_hand dtw	331002 non-null	float64
12	x0_hand pump	331002 non-null	float64
13	x0_handpump	331002 non-null	float64
14	x0_improved spring	331002 non-null	float64
15	x0_insufficient	331002 non-null	float64
16	x0_lake	331002 non-null	float64
17	x0_machine dbh	331002 non-null	float64
18	x0_milky	331002 non-null	float64
19	x0_motorpump	331002 non-null	float64
20	x0_other	331002 non-null	float64
21	x0_rainwater harvesting	331002 non-null	float64
22	x0_river	331002 non-null	float64
23	x0_rope pump	331002 non-null	float64
24	x0_salty	331002 non-null	float64
25	x0_salty abandoned	331002 non-null	float64
26	x0_seasonal	331002 non-null	float64
27	x0_shallow well	331002 non-null	float64
28	x0_soft	331002 non-null	float64
29	x0_spring	331002 non-null	float64
30	x0_submersible	331002 non-null	float64
31	x0_unknown	331002 non-null	float64
32	x0_wind-powered	331002 non-null	float64

```
dtypes: float64(33)
```

memory usage: 93.3 MB

In [23]:

```
1 #Defining numerical df
2 X_train_numerical = X_train[['amount_tsh', 'gps_height']].copy()
3 X_train_numerical.head()
```

Out[23]:

	amount_tsh	gps_height
3607	50.0	2092
50870	0.0	0
20413	0.0	0
52806	0.0	0
50091	300.0	1023

In [24]:

```
1
2 # Initialize MinMaxScaler
3 scaler = MinMaxScaler()
4
5 # Scale the selected column
6 X_train_numerical = scaler.fit_transform(X_train_numerical)
7
8 column_names = ['amount_tsh']
9 # Converting X_train_numerical back into a df
10 X_train_numerical = pd.DataFrame(X_train_numerical, columns=column_names)
11
12 # Display the array
13 X_train_numerical
```

```
-----
NameError                                Traceback (most recent call last)
<ipython-input-24-a113bd572cb9> in <module>
      6
      7 # Converting X_train_numerical back into a df
----> 8 X_train_numerical = pd.DataFrame(X_train_numerical, columns=column_names)
      9
     10 # Display the array
```

NameError: name 'column_names' is not defined

```
In [ ]: 1 X_train_numerical.shape
```

```
In [ ]: 1 # X_train_numerical is a NumPy array
2
3
4
5 # Convert the NumPy array to a pandas DataFrame with specified column names
6 X_train_numerical = pd.DataFrame(X_train_numerical, columns=column_names)
7
8 # Display the DataFrame
9
10
11 print(X_train_numerical)
```

```
In [ ]: 1 X_train_full = pd.concat([X_train_numerical, X_train_categorical], axis=1)
2 X_train_full.head()
```

```
In [ ]: 1 X_train_full.info()
```

```
In [ ]: 1 missing_values = X_train_full.isnull().sum()
2 print(missing_values)
```

```
In [ ]: 1 X_train_full.info()
```

```
In [ ]: 1 # Creating a heatmap from the initial dataframe
2 fig, ax = plt.subplots(figsize=(10,10))
3 cor = Waterwells_df.corr()
4 sns.heatmap(cor, cmap="Blues", annot=True)
```

I wanted to create a function so I could easily evaluate each of models with an r2 score, root mean squared error, and mean absolute error.

```

In [ ]: 1 def evaluate_model(y_test, y_pred, lr):
        2     # R-squared (R2)
        3     r2 = r2_score(y_test, y_pred)
        4
        5     # Root Mean Squared Error (RMSE)
        6     rmse = mean_squared_error(y_test, y_pred, squared=False)
        7
        8     # Mean Absolute Error (MAE)
        9     mae = mean_absolute_error(y_test, y_pred)
       10
       11     # Intercept
       12     #intercept = lr.intercept_
       13
       14     # Printing the results
       15     print("R2 score: ", r2)
       16     print("Root Mean Squared Error: ", rmse)
       17     print("Mean Absolute Error: ", mae)
       18     #print("Intercept: ", intercept)
       19
       20     # Returning the results as a dictionary
       21     results_model = {
       22         'r2': r2,
       23         'rmse': rmse,
       24         'mae': mae,
       25         #'intercept': intercept
       26     }
       27
       28     return results_model

```

Modeling

```

In [ ]: 1 logreg = LogisticRegression(fit_intercept=False, C=1e12, solver='liblinear')
        2 model_log = logreg.fit(X_train_full, y_train)

```

Evaluation

Conclusion

Recommendations

Limitations

Next Steps

In []:

1